



Johnny's Software Lab

Why does my program behave like this?

Tools to help you quickly understand the behavior of
your program



<https://johnysswlab.com>,



@johnysswlab



ivica@johnysswlab.com



About me

- Ivica Bogosavljevic - application performance specialist
- Professional focus is C/C++ application performance improvement:
 - Better algorithms
 - Better exploiting the underlying hardware
 - Better usage of the standard library
 - Better usage of programming language
 - Better usage of the operating system.
- Work as a an external expert
 - If your software is slow, I can help you make it faster
- Writer for software performance blog: Johny's Software Lab - link in the footer
 - For all the people interested in software performance





Introduction

- Starting a work on unfamiliar codebase
 - Hot and cold functions
 - Timeline
 - Memory usage
 - Hardware efficiency
 - How program interacts with the operating system
- Debugging not covered here
 - Debugging weird behavior
 - Debugging memory issues
 - Debugging thread races
 - Difficult to reproduce bugs

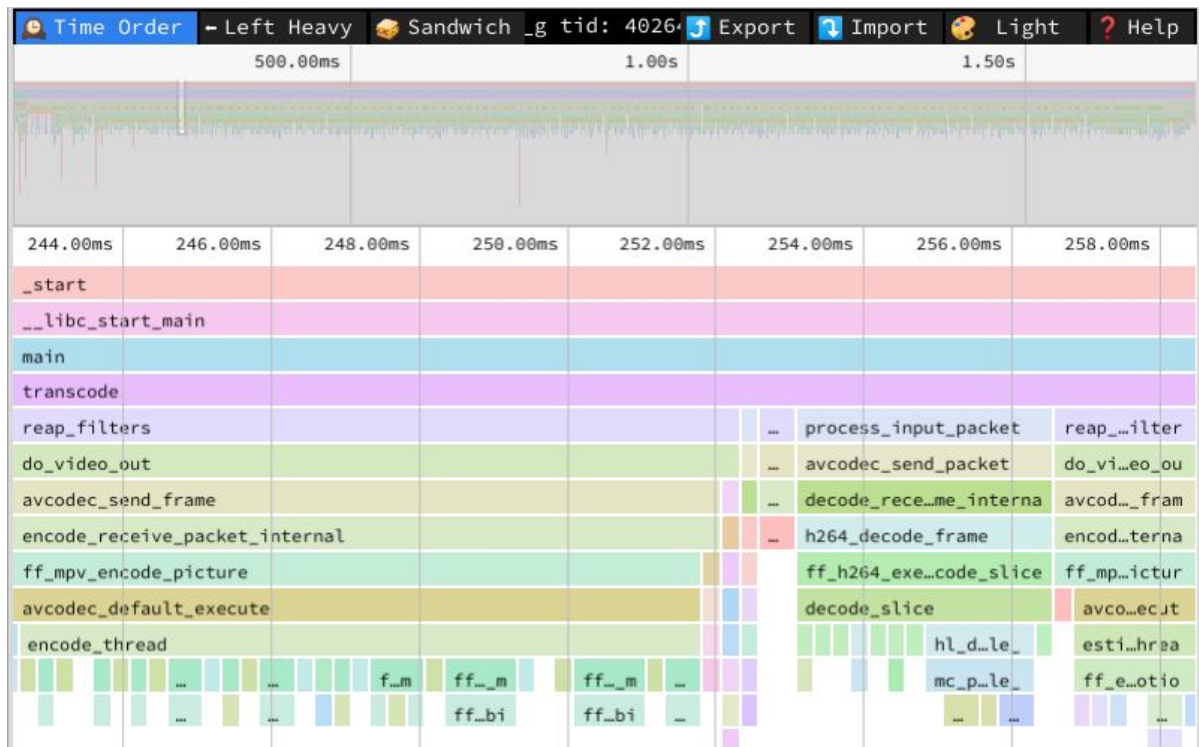
- All examples are available here:

<https://github.com/ibogosavljevic/johnysswlab/tree/master/talks/tools>





Flamegraphs - e.g. speedscope



- FFMPEG Example:
<https://www.speedscope.app/#profileURL=https%3A%2F%2Fraw.githubusercontent.com%2Fibogoslavjevic%2Fjohnnysswlab%2Fmaster%2F2021-03-speedscope%2Fspeedscope-ffmpeg.txt>





Speedscope

- Hands-on in terminal
- Reference
 - ``perf record --call-graph dwarf -F 99 ./my_app``
 - Parameter -F is the sampling frequency, increase for short running processes, decrease for long running processes
 - The app needs to be compiled with debug symbols (``-g`` for GCC and CLANG)
 - `perf script | speedscope -`
- More information about supported platforms and installation:
<https://www.speedscope.app/>
- More information on how to collect data using perf:
- <https://johnysswlab.com/speedscope-visualize-what-your-program-is-doing-and-where-it-is-spending-time/>





Speedscope

speedscope is designed to ingest profiles from a variety of different profilers for different programming languages & environments. Click the links below for documentation on how to import from a specific source.

- JavaScript
 - [Importing from Chrome](#)
 - [Importing from Firefox](#)
 - [Importing from Safari](#)
 - [Importing from Node.js](#)
- Ruby
 - [Importing from stackprof](#)
 - [Importing from rbspy](#)
 - [Importing from ruby-prof](#)
- Python
 - [Importing from py-spy](#)
 - [pyspeedscope](#)
 - [Importing from Austin](#)
- Go
 - [Importing from pprof](#)
- Rust
 - [flamescope](#)
- Native code
 - [Importing from Instruments.app](#) (macOS)
 - [Importing from perf](#) (linux)
- [Importing from .NET Core](#)
- [Importing from GHC \(Haskell\)](#)
- [Importing from custom sources](#)







Heap profilers, e.g. Heaptrack



Flamecharts



Cumulated allocations



Sizes



Consumed

- Debugging memory consumption issues





Heaptrack

- Hands-on in terminal
- Installation through packet manager
- Reference
 - Compile with debug symbols (`-g` on GCC and CLANG)
 - Recording information: `heaptrack ./my_app`
 - Visualization: `heaptrack --analyze heaptrack*.gz`
- Additional information
 - <https://www.youtube.com/watch?v=ZXTI5iWHhrg>
 - https://github.com/adesitter/accu_presentations/blob/master/ReducingMemoryAllocations_CppOnSea_2020.pdf





Coverage tools, e.g. Icov

```
1778      :      /* phixi      */
1779 25164000 :      Real_t norm = Real_t(1.) / (domain.delv_xi(ielem)+ ptiny ) ;
1780      :
1781 25164000 :      switch (bcMask & XI_M) {
1782 24325200 :          case XI_M_COMM: /* needs comm data */
1783 24325200 :              delvm = domain.delv_xi(domain.lxim(ielem)); break ;
1784 838800   :              case XI_M_SYMM: delvm = domain.delv_xi(ielem) ;          break ;
1785      :              case XI_M_FREE: delvm = Real_t(0.0) ;          break ;
1786 0 :              default:      fprintf(stderr, "Error in switch at %s line %d\n",
1787      :                          FILE , LINE );
1788 0 :              delvm = 0; /* ERROR - but quiets the compiler */
1789 0 :              break;
1790      :      }
1791 25164000 :      switch (bcMask & XI_P) {
1792 24325200 :          case XI_P_COMM: /* needs comm data */
1793 24325200 :              delvp = domain.delv_xi(domain.lxip(ielem)) ; break ;
1794 0 :              case XI_P_SYMM: delvp = domain.delv_xi(ielem) ;          break ;
1795      :              case XI_P_FREE: delvp = Real_t(0.0) ;          break ;
1796 0 :              default:      fprintf(stderr, "Error in switch at %s line %d\n",
1797      :                          FILE , LINE );
1798 0 :              delvp = 0; /* ERROR - but quiets the compiler */
1799 0 :              break;
1800      :      }
1801      :
```





Icov

- Hands-on in terminal
- Installation from the repositories
- Reference
 - Compile and link with coverage enabled (`--coverage` and -g` with GCC and CLANG)`
 - Run the binary
 - ``lcov --capture --directory my_dir --output-file coverage.info``
 - ``genhtml coverage.info --output-directory out``
 - ``google-chrome out/index.html``
- More info: <http://ltp.sourceforge.net/coverage/lcov.php>





Kernel call tracing, e.g. strace

% time	seconds	usecs/call	calls	errors	syscall
80,11	0,392332	0	510738		futex
19,67	0,096351	9	10255		brk
0,17	0,000817	21	38		munmap
0,05	0,000242	3	73		mmap
0,00	0,000000	0	7		read
0,00	0,000000	0	24		write
0,00	0,000000	0	9		close
0,00	0,000000	0	16	15	stat
0,00	0,000000	0	10		fstat
0,00	0,000000	0	10		mprotect
0,00	0,000000	0	2		rt_sigaction
0,00	0,000000	0	1		rt_sigprocmask
0,00	0,000000	0	8		pread64
0,00	0,000000	0	1	1	access
0,00	0,000000	0	1		execve
0,00	0,000000	0	2	1	arch_prctl
0,00	0,000000	0	1		sched_getaffinity
0,00	0,000000	0	2		getdents64
0,00	0,000000	0	1		set_tid_address
0,00	0,000000	0	87	78	openat
0,00	0,000000	0	1		set_robust_list
0,00	0,000000	0	1		prlimit64
100,00	0,489742		521288	95	total

- You can deduce a lot of information about program by observing how it talks with the OS
- Also important if the program is spending a lot of time in the system mode





strace

- Hands on in the terminal
- Installed on most Linux'
- Reference
 - ``strace ./my_app`` prints all system calls
 - ``strace -k ./my_app`` prints all system calls with complete stack trace
 - ``strace -c ./my_app`` statistical information about system calls, useful for debugging performance issues
- More info:
<https://johnnysswlab.com/lessons-in-debugging-observe-how-programs-interact-with-the-linux-kernel-with-strace/>
- For girls: <https://jvns.ca/strace-zine-v2.pdf>





Hardware counters and event counters - perf

```
Performance counter stats for './matmul2':
```

949,61 msec	task-clock	#	1,000 CPUs utilized
1	context-switches	#	1,053 /sec
0	cpu-migrations	#	0,000 /sec
8572	page-faults	#	9,027 K/sec
3133973668	cycles	#	3,300 GHz
7106556022	instructions	#	2,27 insn per cycle
885515515	branches	#	932,503 M/sec
1474140	branch-misses	#	0,17% of all branches

```
0,949876045 seconds time elapsed
```

```
0,933857000 seconds user
```

```
0,016031000 seconds sys
```

- Perf is a very powerful tool, here we limit ourselves only to the most basic information





perf

- Hands on in terminal
- Reference
 - ``perf stat ./my program`` - lists the basic events
 - User and system - information about time spent in user and system mode
 - Cycles - Information about spent cycles, corresponding to execution time
 - Instruction - Information about executed instructions
 - Instruction per cycle - ideally 4, but very rarely seen that number
 - List all events ``perf list``
- In general: efficient instructions result in smaller cycle count, but increasing instructions increases cycle count
- For measuring on a piece of code instead of the whole program, see LIKWID:
<https://johnysswlab.com/hardware-performance-counters-the-easy-way-quickstart-likwid-perfctr/>





Join the community!

- cppserbia.slack.com
 - Goran Arandžević
 - Ivica Bogosavljević



<https://johnysswlab.com>,



[@johnysswlab](https://twitter.com/johnysswlab)



ivica@johnysswlab.com