

# Introduction to borrow checking in C++

Serbian C++ User Group meetup,  
Belgrade, August 2023, Beogradska

**Goran Arandjelović**

goran.arandjelovic@gmail.com

goran@cppserbia.com

# Rust Ownership, References and Borrowing

# Ownership

- Each value in Rust has an owner
- There can only be one owner at a time
- When the owner goes out of scope, the value will be dropped

# References

- `&T` -> non-mutable
- `&mut T` -> mutable

# References

```
fn main() {  
    let s1 = String::from("hello");  
  
    let len = calculate_length(&s1);  
  
    println!("The length of '{}' is {}.", s1, len);  
}  
  
fn calculate_length(s: &String) -> usize {  
    s.len()  
}
```

# References

```
fn main() {  
    let mut s = String::from("hello");  
  
    change(&mut s);  
}  
  
fn change(some_string: &mut String) {  
    some_string.push_str(", world");  
}
```

# References

- At any given time, you can have either one mutable reference or any number of immutable references
- References must always be valid

# Borrowing

- Action of creating a reference is called *borrowing*



# Borrowing

```
fn main() {  
    let s = String::from("hello");  
  
    change(&s);  
}  
  
fn change(some_string: &String) {  
    some_string.push_str(", world");  
}
```

# Borrowing

```
fn main() {  
    let mut s = String::from("hello");  
  
    change(&mut s);  
}  
  
fn change(some_string: &mut String) {  
    some_string.push_str(", world");  
}
```

# Borrowing

```
let mut s = String::from("hello");
```

```
let r1 = &mut s;
```

```
let r2 = &mut s; // error
```

```
println!("{}", {}, r1, r2);
```

# Borrowing

```
let mut s = String::from("hello");
```

```
let r1 = &s; // ok
```

```
let r2 = &s; // ok
```

```
let r3 = &mut s; // error
```

```
println!("{}", {}, and {}", r1, r2, r3);
```

# Data race

A data race happens when these three behaviors occur:

- Two or more pointers access the same data at the same time
- At least one of the pointers is being used to write to the data
- There's no mechanism being used to synchronize access to the data

# Remember References?

- At any given time, you can have either one mutable reference or any number of immutable references (no **data race**)
- References must always be valid (no **use-after-free**)

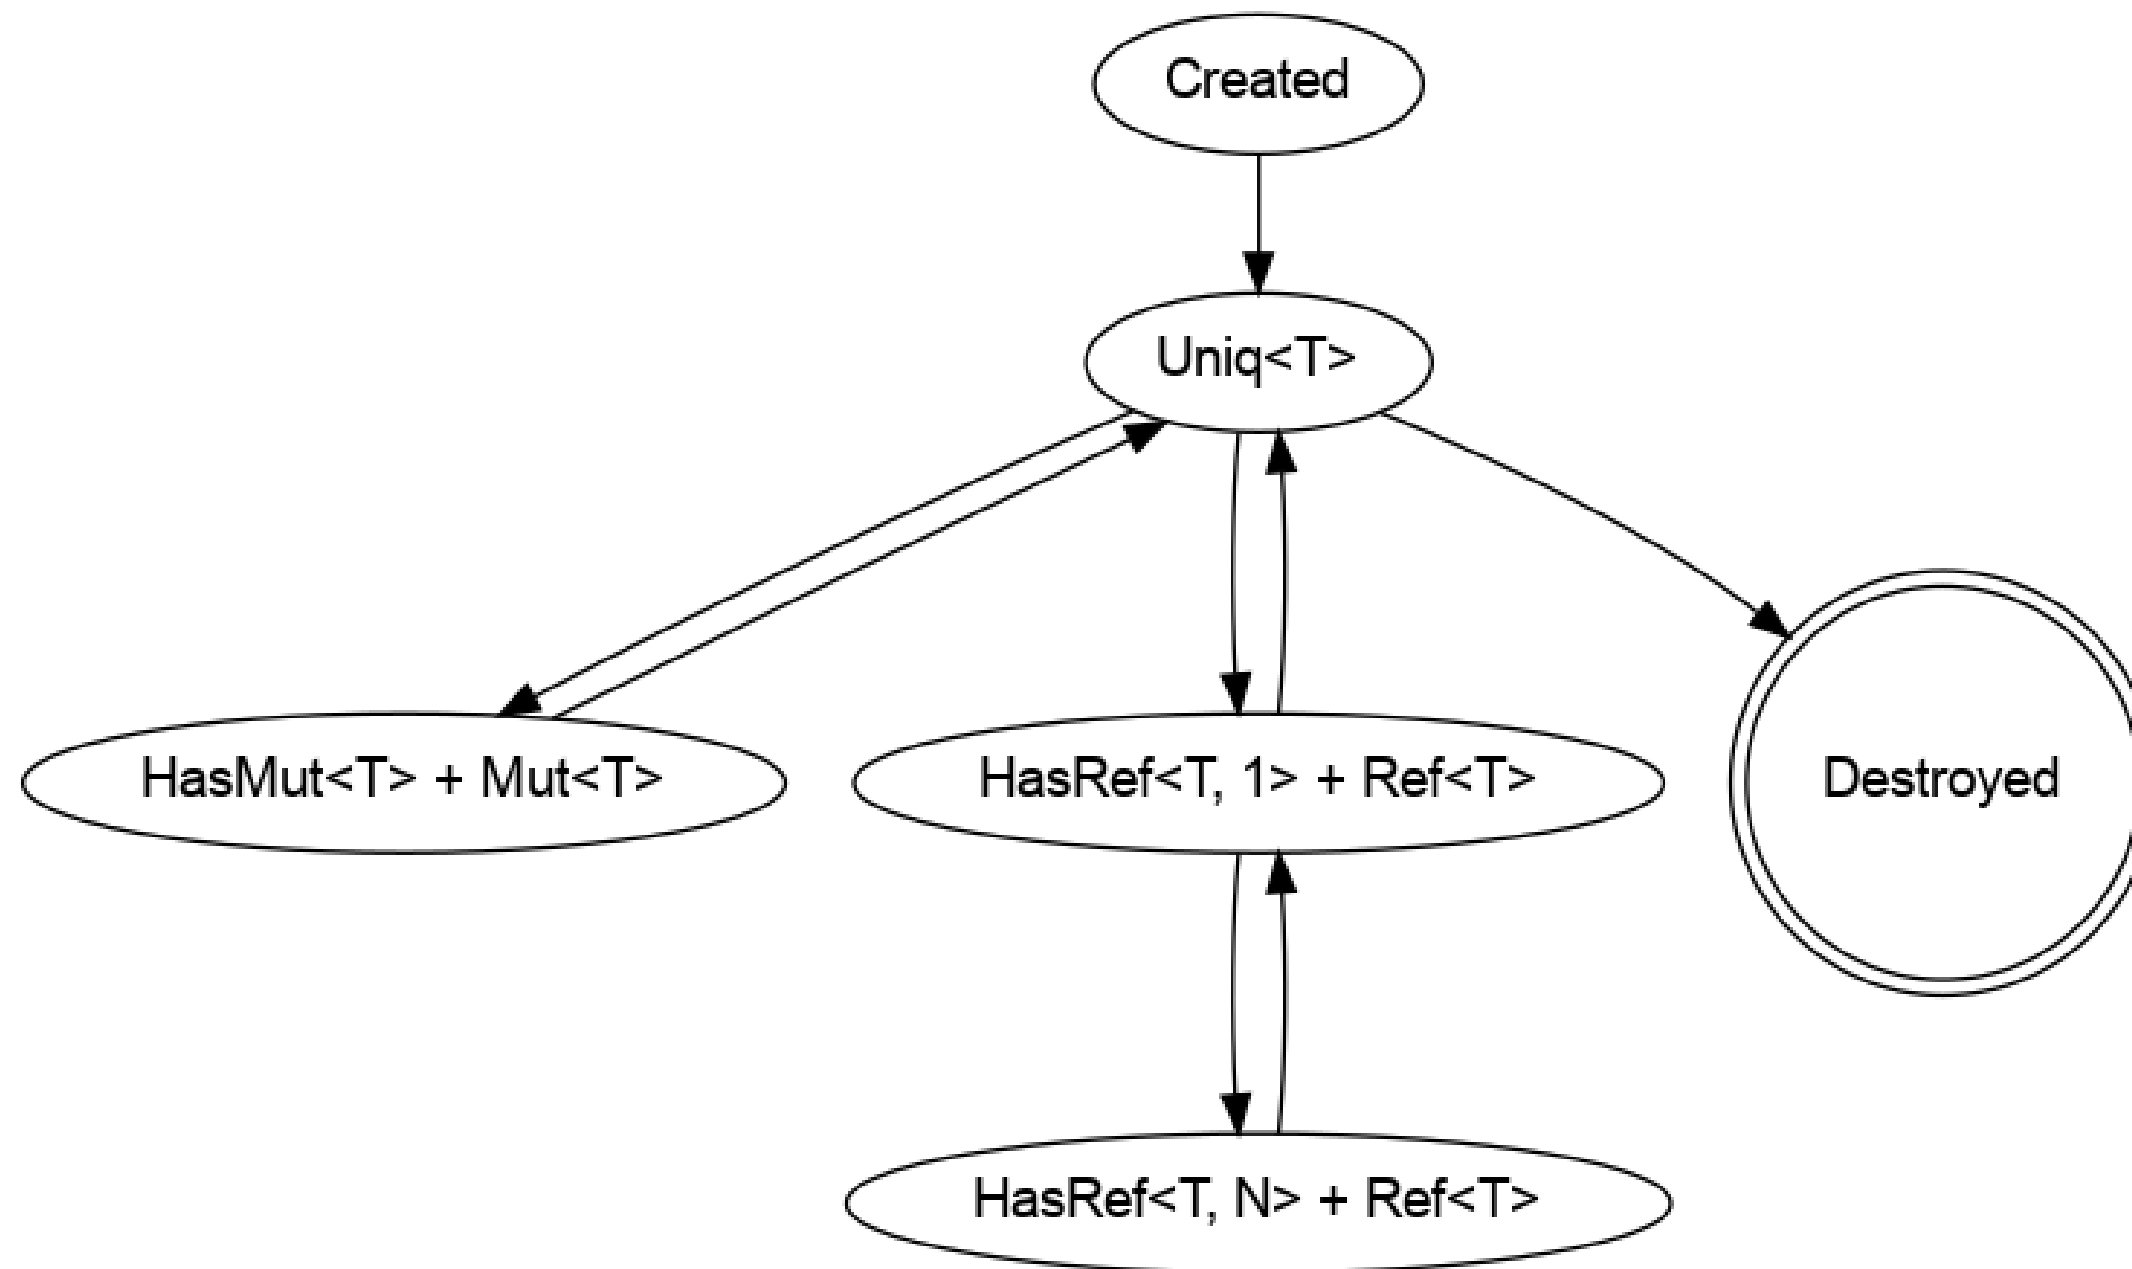
# Borrow checking in C++

# Borrow checking in C++

- Lifetime safety - [wg21.link/p1179](http://wg21.link/p1179)
- Google's Chromium team analyzed if it's possible to implement borrow checker in C++
- Various independent implementations of (runtime) borrow checking



# Borrow checking in C++ — State transition



Author: [palmer@chromium.org](mailto:palmer@chromium.org)

# Borrow checking in C++ — Corresponding types

Define each state as a type:

```
class Uniq<T> {};  
class HasMut<T> {};  
class HasRef<T, unsigned> {};
```

# Borrow checking in C++ — Corresponding types

References require an explicit lifetime, so they also are defined as types:

```
class MutRef<T> {};  
class Ref<T> {};
```

# Borrow checking in C++ — Corresponding types

State transitions free functions from a uniquely owned state to the corresponding references:

```
std::pair<HasMut<T>, MutRef<T>> mut(Uniq<T>);  
std::pair<HasRef<T, 1>, Ref<T>> ref(Uniq<T>);
```

# Borrow checking in C++ — Corresponding types

State transitions free functions back to unique ownership:

```
Uniq<T> consume(HasMut<T>, MutRef<T>);  
Uniq<T> consume(HasRef<T, 1>, Ref<T>);
```

Ignoring HasRef<T, N> where N > 1 for simplicity.

# Borrow checking in C++ — The failures

- Destructors
- Non-destructive moves
- Left-over states

# Borrow checking in C++ — The failures

## Destructors

Object must always be returned to the `Uniq<T>` state/type before it is destroyed. To do otherwise violates the guarantees of the borrow checker. **To do that, `HasMut<T>` and `HasRef<T>` can not be (publicly) destructible.**

- If they are destructible but **do not destroy** the underlying `T`, then we can leak the object beyond the lifetime of the references and unique ownership.
- If they are destructible and **do destroy** the underlying `T`, then we destroy `T` while there is an outstanding reference to it, which creates a UAF scenario.

By ensuring the types `HasMut<T>` and `HasRef<T>` are not (publicly) destructible, we would ensure that they are converted back to `Uniq<T>` eventually, or their usage would fail to compile.

# Borrow checking in C++ — The failures

## Non-destructive moves

- C++ passes ownership through moves, however C++ requires that the destructor of the moved-from object is still run when the object leaves scope
- Rust moves objects destructively. When ownership (always of a uniquely owned object) is transferred, the destructor on the moved-from object would not run, the contents of the object are mem-copied into the destination object and the moved-from object is considered to no longer exist to the compiler
- Each state transition in C++ borrow checker implementation would leave behind a moved-from object representing the old state, which must be destroyed. Because the moved-from object lives outside of the state transition functions, **it requires a public destructor**. This contradicts borrow checker's desired guarantees



# Borrow checking in C++ — The failures

## Left-over states

- Coming back from `HasRef<T>` to `Uniq<T>` creates new (different) `Uniq<T>` object
- What about original `Uniq<T>`? It's a left-over state and we must be sure that nobody will use it for some other state transition (which would unfortunately require runtime check)
- Therefore, going from `Uniq<T>` to `HasRef<T>` **must destroy** `Uniq<T>` state
- If we choose not to destroy it and leave it assignable, we generate another problem: Some other object could be returned into that state, because `Uniq<T>` is bound to a particular type `T`, not the object itself

# Borrow checking in C++ — The failures

## Reference lifetime problem

- Rust compiler guarantees references lifetimes. If we pass/borrow mutable reference to a function, Rust ensures that the reference will be returned back.
- In order to mimic the same lifetime rules for references, C++ function that receives `MutRef<T>` must also return it to the caller. That means putting additional burden on the API

# Borrow checking in C++ — Runtime

- Runtime borrow checking is possible
- There are a few good prototypes out there (e.g. <https://github.com/shuaimu/borrow-cpp/>)
- Some compilers and tool-chains already support UAF (use-after-free) analysis to some extent
- Using a runtime implementation + static analysis can help us prevent the most common errors

# References (not the Rust ones)

- Lifetime safety - [wg21.link/p1179](http://wg21.link/p1179)
- [Borrowing Trouble: The Difficulties Of A C++ Borrow-Checker](#)
- [Destructive move | N4034](#)
- [Sean Parent - About move](#)
- [Rust - References and Borrowing](#)
- [Bjørn Reese - Almost Affine](#)
- [Example of stateful metaprogramming](#)

Thank you